

# Examining the State-of-the-Art in News Timeline Summarization

Demian Gholipour Ghalandari, Georgiana Ifrim

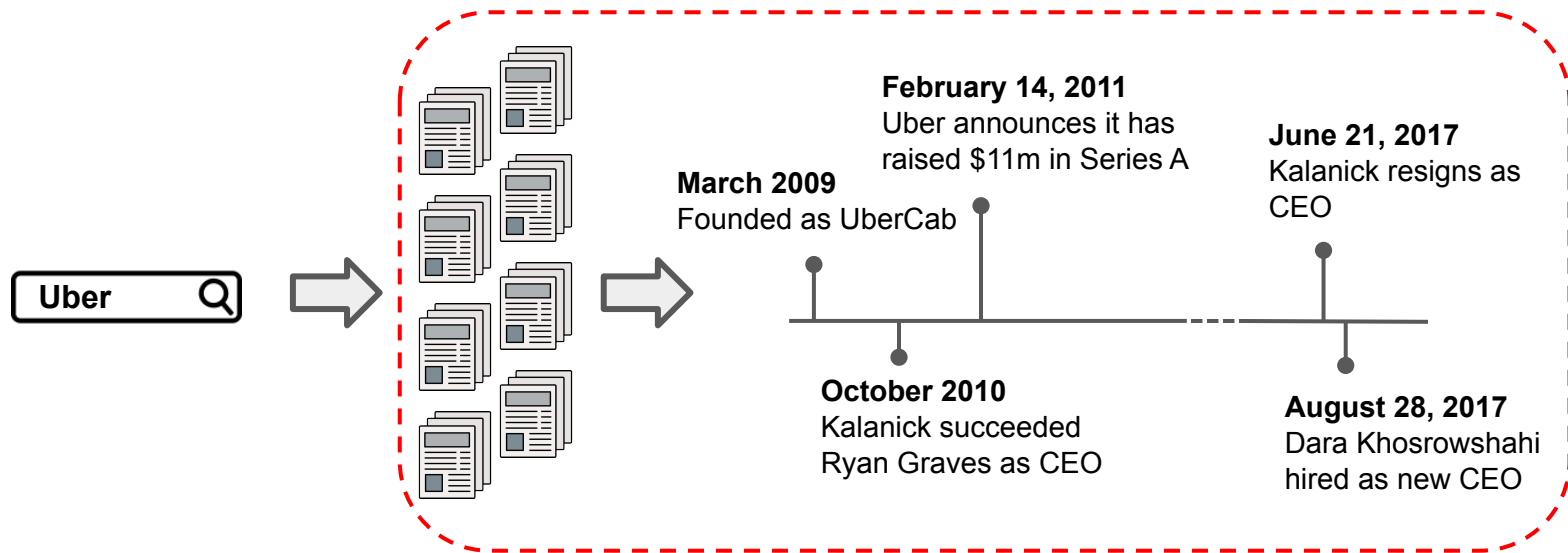
School of Computer Science  
University College Dublin, Ireland



# Outline

- Timeline Summarization (TLS) Task
- Previous Work
- Types of Approaches
  - Date-wise
  - Clustering-based
- New Dataset
- Experiments
- Conclusions

# What is Timeline Summarization (TLS)?



## TASK

- **Input:** Article collection **A** + query (topic) keyphrases **Q**
- **Output:** List of <date, summary> pairs
- **Length constraints:** **L** dates, with **M** total sentences,  
or **K=M/L** sentences per date

# Previous Work in TLS

## Shortcomings in previous work

- Papers about isolated subtasks: date selection, date summarisation
- Inappropriate evaluation metrics, inconsistent experimental settings
- Small ground-truth datasets

## Most recent SOTA (Martschat 2018')


- Establishes experimental settings, uses appropriate evaluation metrics
- Method not scalable
- Poor comparison to previous approaches




Not clear how TLS can be approached and how well different strategies work

# Approaches to Build News Timelines


## Date-wise



## Clustering / Event-based



## One-stage summarization (current SOTA)

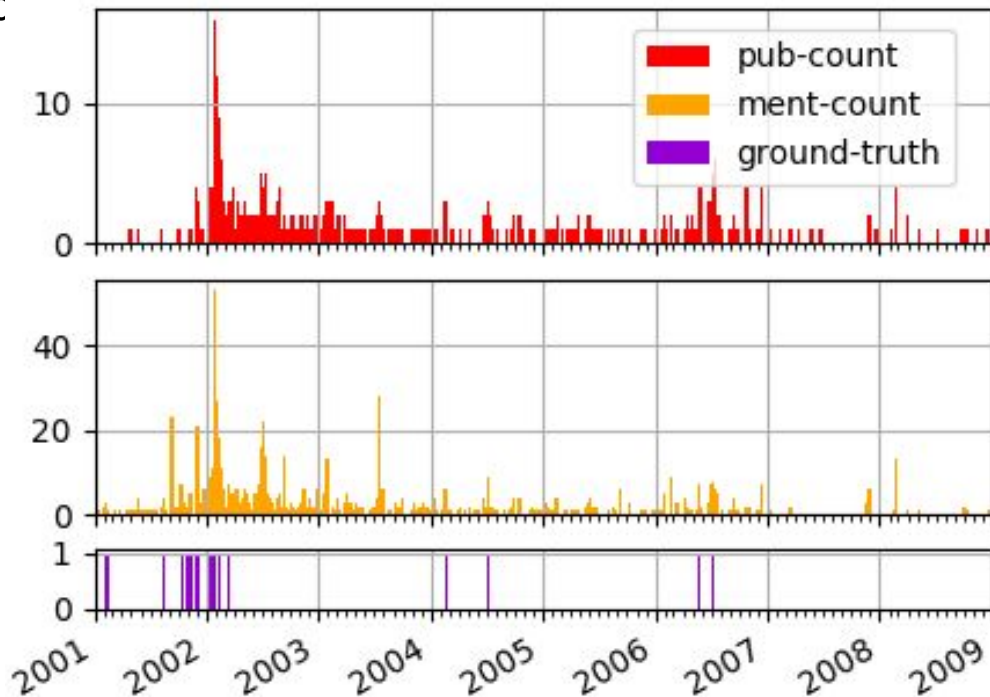


## This work:

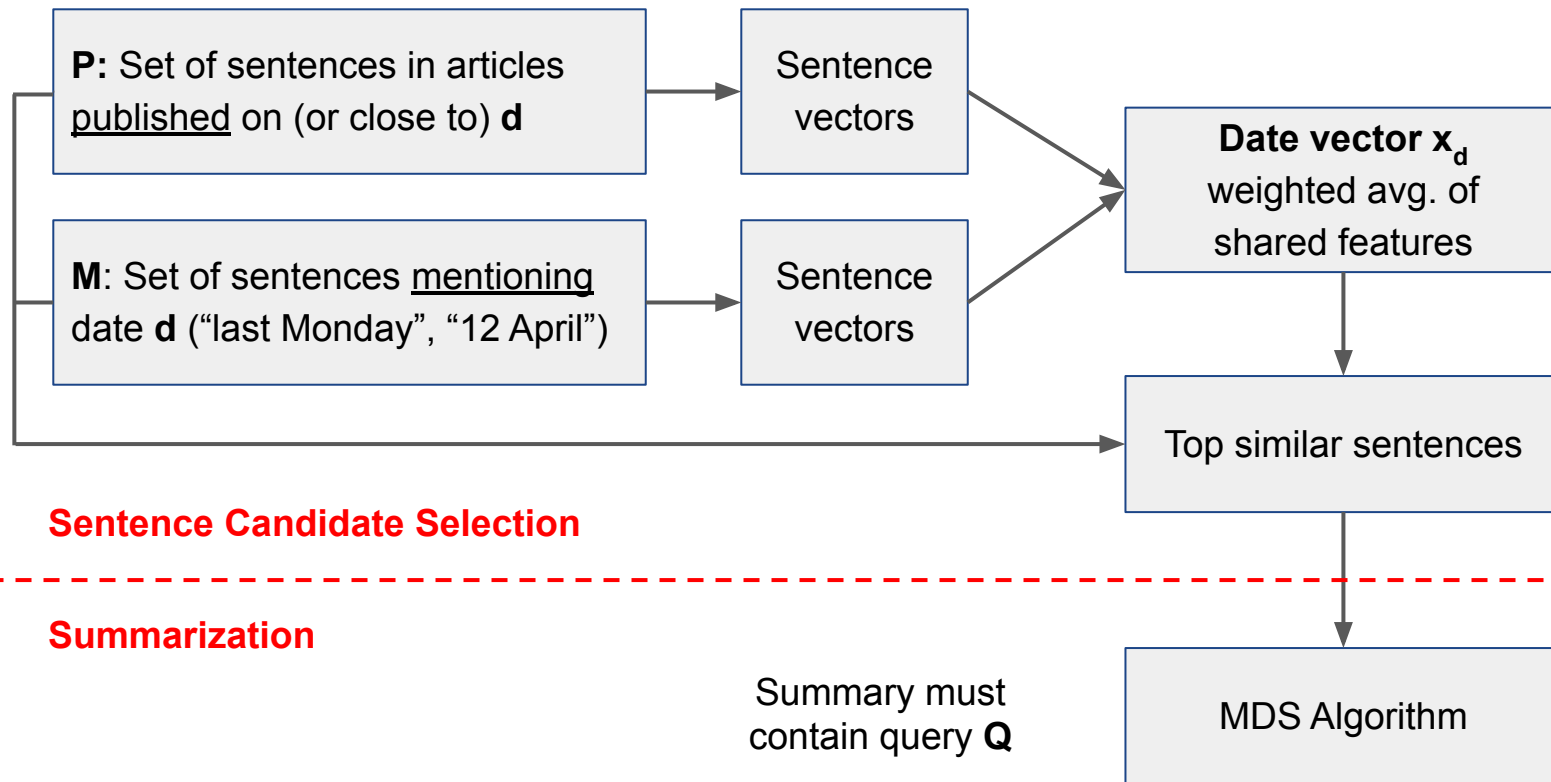
- Implement + evaluate date-wise & clustering-based TLS approaches
- Side-by-side comparison
- New & efficient SOTA by date-wise approach

## Date-wise Approach (1/2): Date Selection

- Task: select  $L$  important dates that should appear on the timeline
- Regression with date features (Tran et al. 2013)
  - # published articles
  - # date is mentioned in text (e.g. 'yesterday', '2 May')
    - before/after/same day
- Alternative approaches
  - pub-count only
  - mention-count only
  - graph-based



## Date-wise Approach (2/2): Summarizing a Particular Date



# Clustering-based Approach

Select L important *events* (clusters), summarize these separately

## Clustering

- Markov clustering on temporal article graph
- Nodes: articles
- Edges: cosine similarity between article vectors
- Edge only if publication dates are 1 day apart

## Cluster dates

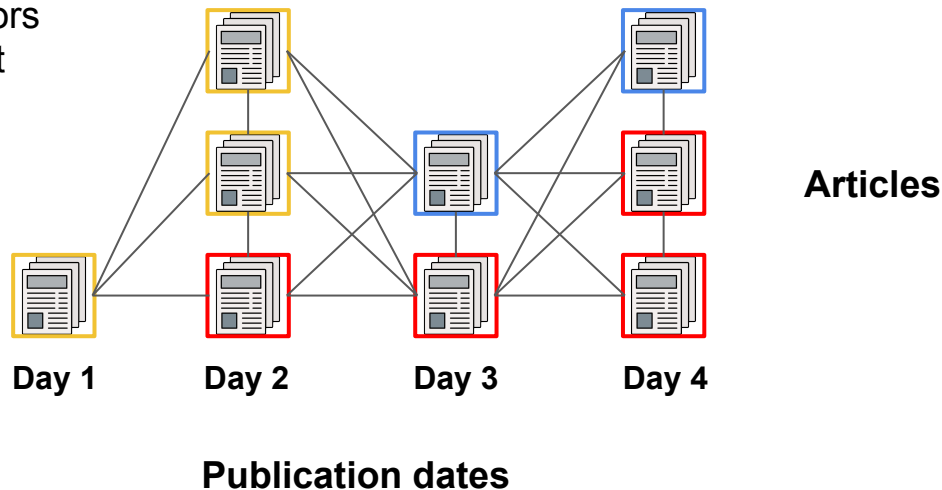
- Most frequently mentioned date in cluster

## Cluster ranking: options

- Cluster size
- **Mention count of cluster date**
- Regression

## Cluster summarization

- Unsupervised MDS algorithm





## New Dataset for TLS (“Entities”)

	#Topics	#TLs	Avg. articles	Avg. summary sentences	Avg. TL-dates	Avg. duration
T17	9	19	508	2.9	36	212 days
Crisis	4	22	2310	1.3	29	343 days
Entities	47	47	959	1.2	23	~12 years

- Larger dataset & longer time periods
- Ground-truth: timelines extracted from “CNN Fast Facts”
- Queries: timeline titles (e.g. “Steve Jobs”)
- Input articles: from TheGuardian API
- Adjustments & Filtering: ensuring overlap between timeline & articles

# Evaluation

- Metrics (Martschat '17)
  - Alignment-based ROUGE (AR) F-score
  - Date Selection F-score
- Leave-one-out cross-validation over individual tasks in a dataset

	T17 Dataset			Crisis Dataset			Entities Dataset		
	AR1-F	AR2-F	Date-F	AR1-F	AR2-F	Date-F	AR1-F	AR2-F	Date-F
<b>One-stage (Martschat'18)</b>	0.105	0.03	0.544	0.075	0.016	0.281	0.042	0.009	0.167
<b>Date-wise</b>	0.12	0.035	0.544	0.089	0.026	0.295	0.057	0.017	0.205
<b>Clustering-based</b>	0.082	0.02	0.407	0.061	0.013	0.226	0.051	0.015	0.174
<b>Oracle</b>	0.312	0.128	0.926	0.367	0.15	0.974	0.232	0.075	0.757

# Observations & Discussion

## Redundancy

- Common phenomenon: adjacent dates in automatic timelines (avoided by clustering-based TLS)
- Forcing date-wise approach to avoid redundancy decreases performance!

**Clustering-based approach:** often gets granularity wrong

## Speed

- Date-wise & clustering-based much faster than submodular one-stage approach

## When is the task particularly difficult?

- High compression ratio: large article collection, short timelines

## Ground-truth

- Subjective in picking important events & formulating summaries (editorial decisions)

## Conclusions & Future Outlook

- Compared 3 TLS strategies (date-wise, clustering, one-stage)
  - Date-wise approach obtains new SOTA
  - Date-wise & clustering very efficient compared to one-stage
- Textual date mentions are very helpful
- TLS is a hard task, difficult to imitate ground-truth
- Redundancy remains a problem
- Quality: reasonable apart from occasional redundancy & style issues

What can be improved?

- Better datasets: larger, more news sources, more timelines
- End-to-end methods, abstractive techniques
- Alternative automatic evaluation, human evaluation